

mnist-ocr-cnn-guntas

May 12, 2023


```
[1]: from tensorflow.keras.datasets import mnist
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.optimizers import SGD
from sklearn.preprocessing import LabelBinarizer
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report
from imutils import build_montages
import matplotlib.pyplot as plt
import numpy as np
import argparse
import cv2

((trainData, trainLabels), (testData, testLabels)) = mnist.load_data()
data = np.vstack([trainData, testData])
labels = np.hstack([trainLabels, testLabels])
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>
11490434/11490434 [=====] - 0s 0us/step

```
[2]: fig, axes = plt.subplots(10, 10, figsize = (8, 8), gridspec_kw = dict(hspace = 0.1,
    ↳wspace = 0.1), subplot_kw = {"xticks" : [], "yticks" : []})

for i, ax in enumerate(axes.flat):
    ax.imshow(data[i], cmap = "binary", interpolation = "nearest")
    ax.text(0.05, 0.05, str(labels[i]), transform = ax.transAxes, color = 0.5,
    ↳"green")
```



```
[ ]: dataFlat = data.reshape(data.shape[0], -1)
dataFlat.shape
```

```
[ ]: (70000, 784)
```

```
[ ]: from sklearn.cluster import KMeans
from sklearn.metrics import accuracy_score
from scipy.stats import mode
from warnings import simplefilter
simplefilter(action = "ignore", category = FutureWarning)

km = KMeans(n_clusters = 10, n_init = "auto")
km.fit(dataFlat)
```

```

pred_labels = km.predict(dataFlat)

lab = np.zeros_like(pred_labels)
for i in range(10):
    mask = (pred_labels == i)
    lab[mask] = mode(labels[mask])[0]

print(lab)
print(labels)
print(accuracy_score(lab, labels))

```

```

[8 0 4 ... 7 8 6]
[5 0 4 ... 4 5 6]
0.5807

```

```

[ ]: from sklearn.naive_bayes import GaussianNB
GNB = GaussianNB()
Xtrain, Xtest, Ytrain, Ytest = train_test_split(dataFlat, labels, train_size = 0.7)
GNB.fit(Xtrain, Ytrain)
y_predict = GNB.predict(Xtest)

print(y_predict)
print(Ytest)
print(accuracy_score(y_predict, Ytest))

```

```

[1 0 6 ... 9 6 8]
[1 4 6 ... 7 8 4]
0.5615238095238095

```

```

[5]: print(trainData.shape)
print(trainLabels.shape)
print(testData.shape)
print(testLabels.shape)

```

```

(60000, 28, 28)
(60000,)
(10000, 28, 28)
(10000,)

```

```

[8]: Xtrain = trainData.reshape(trainData.shape[0], trainData.shape[1], trainData.
    ↪shape[2], 1)
Xtest = testData.reshape(testData.shape[0], testData.shape[1], testData.
    ↪shape[2], 1)
print(Xtrain.shape)

```

```

(60000, 28, 28, 1)

```

```
[11]: print(trainData[0])
```


```
[[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0  3  18  18  18 126 136
 175 26 166 255 247 127 0 0 0 0]
 [ 0  0  0  0  0  0  0  0 30 36 94 154 170 253 253 253 253 253
 225 172 253 242 195 64 0 0 0 0]
 [ 0  0  0  0  0  0  0 49 238 253 253 253 253 253 253 253 253 251
 93 82 82 56 39 0 0 0 0 0]
 [ 0  0  0  0  0  0  0 18 219 253 253 253 253 253 198 182 247 241
 0 0 0 0 0 0 0 0 0 0]
 [ 0  0  0  0  0  0  0  0 80 156 107 253 253 205 11 0 43 154
 0 0 0 0 0 0 0 0 0 0]
 [ 0  0  0  0  0  0  0  0  0 14 1 154 253 90 0 0 0 0
 0 0 0 0 0 0 0 0 0 0]
 [ 0  0  0  0  0  0  0  0  0  0  0 139 253 190 2 0 0 0
 0 0 0 0 0 0 0 0 0 0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0 11 190 253 70 0 0
 0 0 0 0 0 0 0 0 0 0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0 35 241 225 160 108 1
 0 0 0 0 0 0 0 0 0 0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0 81 240 253 253 119
 25 0 0 0 0 0 0 0 0 0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0 45 186 253 253
150 27 0 0 0 0 0 0 0 0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0 0 16 93 252
253 187 0 0 0 0 0 0 0 0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0 0 0 0 249
253 249 64 0 0 0 0 0 0 0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0 46 130 183 253
253 207 2 0 0 0 0 0 0 0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0 39 148 229 253 253 253
250 182 0 0 0 0 0 0 0 0]
 [ 0  0  0  0  0  0  0  0  0  0 24 114 221 253 253 253 253 201
78 0 0 0 0 0 0 0 0 0]
 [ 0  0  0  0  0  0  0 23 66 213 253 253 253 253 198 81 2
0 0 0 0 0 0 0 0 0]
 [ 0  0  0  0  0  0 18 171 219 253 253 253 253 195 80 9 0 0
```

```

0 0 0 0 0 0 0 0 0 0]
[ 0 0 0 0 55 172 226 253 253 253 253 244 133 11 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0]
[ 0 0 0 0 136 253 253 253 212 135 132 16 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0]]

```

```
[9]: plt.imshow(Xtrain[0])
plt.show()
```



0.1 Categorization of Labels

For each element in the `trainLabels`, an array would be created in the `Ytrain` with size that of the number of distinct labels and only the entry corresponding to the current label will be 1 rest all 0.

```
[12]: from keras.utils import to_categorical
```

```
Ytrain = to_categorical(trainLabels)
Ytest = to_categorical(testLabels)
print(trainLabels)
print(Ytrain)
```

```
[5 0 4 ... 5 6 8]
[[0. 0. 0. ... 0. 0. 0.]
 [1. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]
 ...
 [0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 1. 0.]]
```

0.2 Building the CNN

- **Conv2D(N, kernel_size = p)** N represents the number of nodes in this layer and p x p is the dimension of the kernel matrix
- relu -> **ReLU : Rectified Linear Activation**
- Activation functions like the **sigmoid, tanh, ReLU(max(x, 0))**, step map the input to a value between some range (usually -1 to 1)
- input_shape -> Shape of each input image
- The **"Flatten()"** connects the Convolutional layers to the Dense layers of the Neural Network
- **Dense(x, activation)** -> x represents the number of possible outputs (0 - 9 in our case)
- **Softmax** function outputs the values with all summing upto 1 like they are probabilities

```
[13]: from keras.models import Sequential
from keras.layers import Dense, Conv2D, Flatten

model = Sequential()

model.add(Conv2D(64, kernel_size = 3, activation = "relu", input_shape = Xtrain.
    ↪shape[1:]))
model.add(Conv2D(32, kernel_size = 3, activation = "relu"))
model.add(Flatten())
model.add(Dense(10, activation = "softmax"))
```

0.3 Compiling the CNN

- We require three parameters for compiling : **optimizer, loss, metrics**.
- **Optimizer** decides the **Learning Rate**, LR determines the Weights in each layer
- We have chose **"adam"** as our optimizer
- **Loss** and **metrics** tell our model on which factor's value should it decide/change the weight in each step

```
[14]: model.compile(optimizer = "adam", loss = "categorical_crossentropy", metrics = [
      ↪ "accuracy"])
```

0.4 Training the Model

```
[15]: model.fit(Xtrain, Ytrain, validation_data = (Xtest, Ytest), epochs = 3)
```

```
Epoch 1/3
1875/1875 [=====] - 212s 112ms/step - loss: 0.1965 -
accuracy: 0.9536 - val_loss: 0.0878 - val_accuracy: 0.9727
Epoch 2/3
1875/1875 [=====] - 229s 122ms/step - loss: 0.0640 -
accuracy: 0.9800 - val_loss: 0.0749 - val_accuracy: 0.9767
Epoch 3/3
1875/1875 [=====] - 259s 138ms/step - loss: 0.0483 -
accuracy: 0.9843 - val_loss: 0.0773 - val_accuracy: 0.9768
```

```
[15]: <keras.callbacks.History at 0x7f391ab1cb50>
```

0.5 Predictions

```
[22]: from sklearn.metrics import accuracy_score
y_model = model.predict(Xtest[:100])
pred = np.argmax(y_model, axis = 1)
actual = np.argmax(Ytest[:100], axis = 1)
print(pred)
print(actual)
print(accuracy_score(pred, actual))
```

```
4/4 [=====] - 0s 31ms/step
[7 2 1 0 4 1 4 9 6 9 0 6 9 0 1 5 9 7 8 4 9 6 6 5 4 0 7 4 0 1 3 1 3 4 7 2 7
 1 2 1 1 7 4 2 3 5 1 2 4 4 6 3 5 5 6 0 4 1 9 5 7 8 9 3 7 4 6 4 3 0 7 0 2 9
 1 7 3 2 9 7 7 6 2 7 8 4 7 3 6 1 3 6 9 3 1 6 1 7 6 9]
[7 2 1 0 4 1 4 9 5 9 0 6 9 0 1 5 9 7 3 4 9 6 6 5 4 0 7 4 0 1 3 1 3 4 7 2 7
 1 2 1 1 7 4 2 3 5 1 2 4 4 6 3 5 5 6 0 4 1 9 5 7 8 9 3 7 4 6 4 3 0 7 0 2 9
 1 7 3 2 9 7 7 6 2 7 8 4 7 3 6 1 3 6 9 3 1 4 1 7 6 9]
0.97
```