

# mnist-cnn

May 12, 2023

## 0.1 Getting the MNIST Dataset

```
[ ]: from tensorflow.keras.datasets import mnist
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.optimizers import SGD
from sklearn.preprocessing import LabelBinarizer
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report
from imutils import build_montages
import matplotlib.pyplot as plt
import numpy as np
import argparse
import cv2


((trainData, trainLabels), (testData, testLabels)) = mnist.load_data()
data = np.vstack([trainData, testData])
labels = np.hstack([trainLabels, testLabels])
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>  
11490434/11490434 [=====] - 0s 0us/step

## 0.2 The MNIST Digits

```
[ ]: fig, axes = plt.subplots(10, 10, figsize = (8, 8), gridspec_kw = dict(hspace = 0.1, wspace = 0.1), subplot_kw = {"xticks" : [], "yticks" : []})

for i, ax in enumerate(axes.flat):
    ax.imshow(data[i], cmap = "binary", interpolation = "nearest")
    ax.text(0.05, 0.05, str(labels[i]), transform = ax.transAxes, color = "green")
```



```
[ ]: dataFlat = data.reshape(data.shape[0], -1)
dataFlat.shape
```

```
[ ]: (70000, 784)
```

### 0.3 KMeans Clustering on the Digits

```
[ ]: from sklearn.cluster import KMeans
from sklearn.metrics import accuracy_score
from scipy.stats import mode
from warnings import simplefilter
simplefilter(action = "ignore", category = FutureWarning)
```

```

km = KMeans(n_clusters = 10, n_init = "auto")
km.fit(dataFlat)
pred_labels = km.predict(dataFlat)

lab = np.zeros_like(pred_labels)
for i in range(10):
    mask = (pred_labels == i)
    lab[mask] = mode(labels[mask])[0]

print(lab)
print(labels)
print(accuracy_score(lab, labels))

```

```

[8 0 4 ... 7 8 6]
[5 0 4 ... 4 5 6]
0.5807

```

## 0.4 Naive Bayes Classifier on the Digits

```

[ ]: from sklearn.naive_bayes import GaussianNB
GNB = GaussianNB()
Xtrain, Xtest, Ytrain, Ytest = train_test_split(dataFlat, labels, train_size = 0.7)
GNB.fit(Xtrain, Ytrain)
y_predict = GNB.predict(Xtest)

print(y_predict)
print(Ytest)
print(accuracy_score(y_predict, Ytest))

```

```

[1 0 6 ... 9 6 8]
[1 4 6 ... 7 8 4]
0.5615238095238095

```

## 0.5 Preparing the Data for the CNN

```

[ ]: print(trainData.shape)
print(trainLabels.shape)
print(testData.shape)
print(testLabels.shape)

```

```

(60000, 28, 28)
(60000,)
(10000, 28, 28)
(10000,)

```

```
[ ]: Xtrain = trainData.reshape(trainData.shape[0], trainData.shape[1], trainData.
    ↪shape[2], 1)
Xtest = testData.reshape(testData.shape[0], testData.shape[1], testData.
    ↪shape[2], 1)
print(Xtrain.shape)
```

(60000, 28, 28, 1)

```
[ ]: print(trainData[0])
```

```
[[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0  3  18  18  18 126 136
 175 26 166 255 247 127 0 0 0 0]
 [ 0  0  0  0  0  0  0  0  30 36 94 154 170 253 253 253 253 253
 225 172 253 242 195 64 0 0 0 0]
 [ 0  0  0  0  0  0  0 49 238 253 253 253 253 253 253 253 253 251
 93 82 82 56 39 0 0 0 0 0]
 [ 0  0  0  0  0  0  0 18 219 253 253 253 253 253 198 182 247 241
 0 0 0 0 0 0 0 0 0 0]
 [ 0  0  0  0  0  0  0  0 80 156 107 253 253 205 11 0 43 154
 0 0 0 0 0 0 0 0 0 0]
 [ 0  0  0  0  0  0  0  0  0 14 1 154 253 90 0 0 0 0
 0 0 0 0 0 0 0 0 0 0]
 [ 0  0  0  0  0  0  0  0  0  0  0 139 253 190 2 0 0 0
 0 0 0 0 0 0 0 0 0 0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0 11 190 253 70 0 0 0
 0 0 0 0 0 0 0 0 0 0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0 35 241 225 160 108 1
 0 0 0 0 0 0 0 0 0 0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0 81 240 253 253 119
 25 0 0 0 0 0 0 0 0 0 0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0 45 186 253 253
 150 27 0 0 0 0 0 0 0 0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0 0 16 93 252
 253 187 0 0 0 0 0 0 0 0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0 0 0 0 249
 253 249 64 0 0 0 0 0 0 0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0 46 130 183 253
```

```


253 207 2 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 0 0 39 148 229 253 253 253
250 182 0 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 0 24 114 221 253 253 253 253 201
78 0 0 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 23 66 213 253 253 253 253 198 81 2
0 0 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 18 171 219 253 253 253 253 195 80 9 0 0
0 0 0 0 0 0 0 0 0]
[ 0 0 0 0 55 172 226 253 253 253 253 244 133 11 0 0 0 0
0 0 0 0 0 0 0 0]
[ 0 0 0 0 136 253 253 253 212 135 132 16 0 0 0 0 0 0
0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0]

```

```

[ ]: plt.imshow(Xtrain[0])
plt.show()

```



## 0.6 Categorization of Labels

For each element in the trainLabels, an array would be created in the Ytrain with size that of the number of distinct labels and only the entry corresponding to the current label will be 1 rest all 0.

```
[ ]: from keras.utils import to_categorical
```

```
Ytrain = to_categorical(trainLabels)
Ytest = to_categorical(testLabels)
print(trainLabels)
print(Ytrain)
```

```
[5 0 4 ... 5 6 8]
[[0. 0. 0. ... 0. 0. 0.]
 [1. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]
 ...
 [0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 1. 0.]]
```

## 0.7 Building the CNN

- **Conv2D(N, kernel\_size = p)** N represents the number of nodes in this layer and p x p is the dimension of the kernel matrix
- relu -> **ReLU : Rectified Linear Activation**
- Activation functions like the **sigmoid, tanh, ReLU(max(x, 0)), step** map the input to a value between some range (usually -1 to 1)
- input\_shape -> Shape of each input image
- The “**Flatten()**” connects the Convolutional layers to the Dense layers of the Neural Network
- **Dense(x, activation)** -> x represents the number of possible outputs (0 - 9 in our case)
- **Softmax** function outputs the values with all summing upto 1 like they are probabilities

```
[ ]: from keras.models import Sequential
from keras.layers import Dense, Conv2D, Flatten

model = Sequential()

model.add(Conv2D(64, kernel_size = 3, activation = "relu", input_shape = Xtrain.
    ↪shape[1:]))
model.add(Conv2D(32, kernel_size = 3, activation = "relu"))
model.add(Flatten())
model.add(Dense(10, activation = "softmax"))
```

## 0.8 Compiling the CNN

- We require three parameters for compiling : **optimizer, loss, metrics**.
- **Optimizer** decides the **Learning Rate**, LR determines the Weights in each layer
- We have chose “**adam**” as our optimizer

- **Loss** and **metrics** tell our model on which factor's value should it decide/change the weight in each step

```
[ ]: model.compile(optimizer = "adam", loss = "categorical_crossentropy", metrics = [
    ↪ "accuracy"])
```

## 0.9 Training the Model

```
[ ]: model.fit(Xtrain, Ytrain, validation_data = (Xtest, Ytest), epochs = 3)
```

```
Epoch 1/3
1875/1875 [=====] - 212s 112ms/step - loss: 0.1965 -
accuracy: 0.9536 - val_loss: 0.0878 - val_accuracy: 0.9727
Epoch 2/3
1875/1875 [=====] - 229s 122ms/step - loss: 0.0640 -
accuracy: 0.9800 - val_loss: 0.0749 - val_accuracy: 0.9767
Epoch 3/3
1875/1875 [=====] - 259s 138ms/step - loss: 0.0483 -
accuracy: 0.9843 - val_loss: 0.0773 - val_accuracy: 0.9768
```

```
[ ]: <keras.callbacks.History at 0x7f391ab1cb50>
```

## 0.10 Predictions

```
[ ]: from sklearn.metrics import accuracy_score
y_model = model.predict(Xtest[:100])
pred = np.argmax(y_model, axis = 1)
actual = np.argmax(Ytest[:100], axis = 1)
print(pred)
print(actual)
print(accuracy_score(pred, actual))
```

```
4/4 [=====] - 0s 31ms/step
[7 2 1 0 4 1 4 9 6 9 0 6 9 0 1 5 9 7 8 4 9 6 6 5 4 0 7 4 0 1 3 1 3 4 7 2 7
 1 2 1 1 7 4 2 3 5 1 2 4 4 6 3 5 5 6 0 4 1 9 5 7 8 9 3 7 4 6 4 3 0 7 0 2 9
 1 7 3 2 9 7 7 6 2 7 8 4 7 3 6 1 3 6 9 3 1 6 1 7 6 9]
[7 2 1 0 4 1 4 9 5 9 0 6 9 0 1 5 9 7 3 4 9 6 6 5 4 0 7 4 0 1 3 1 3 4 7 2 7
 1 2 1 1 7 4 2 3 5 1 2 4 4 6 3 5 5 6 0 4 1 9 5 7 8 9 3 7 4 6 4 3 0 7 0 2 9
 1 7 3 2 9 7 7 6 2 7 8 4 7 3 6 1 3 6 9 3 1 4 1 7 6 9]
0.97
```